

## Comparative Analysis of Handwritten Alphabet Recognition Models

### 1. Overview

This document presents a comparative analysis of four deep learning architectures for handwritten alphabet (A-Z) recognition using the **A-Z Handwritten Data.csv** dataset (28x28 grayscale images, 26 classes).

### 2. Model Performance Summary

Model	Test Accuracy	Precision	Recall	F1-Score	AUC (Avg)	Parameters
<b>MobileNet (Custom)</b>	<b>99.03%</b>	<b>98.66%</b>	<b>98.74%</b>	<b>98.69%</b>	~0.999	<b>3.27M</b>
<b>VGG19 (Scratch)</b>	98.47%	97.88%	97.91%	97.88%	~0.998	3.19M
<b>ResNet50V2</b>	97.06%	96.40%	95.62%	95.98%	~0.997	24.68M
<b>InceptionV1 (GoogLeNet)</b>	95.27%	94.72%	93.25%	93.94%	~0.995	6.20M

### 3. Detailed Model Analysis

#### 3.1 VGG19 (Built from Scratch)

**Architecture:** Classical deep CNN with small 3x3 filters, multiple convolutional blocks, and fully connected layers.

**Pros:**

- Simple, interpretable architecture
- Excellent performance (98.47%) despite being built from scratch
- Minimal preprocessing required (28x28 grayscale input)
- Lower parameter count (3.19M) than pretrained models

**Cons:**

- Less efficient than MobileNet
- Can be prone to overfitting (requires careful regularization)
- No transfer learning benefits

**Why it performed well:**

The dataset (28x28 images) is well-suited for VGG's hierarchical feature extraction. The model depth (16 layers) captures intricate patterns in handwritten characters effectively without being too complex for the task.

#### 3.2 ResNet50V2 (Pretrained)

**Architecture:** Deep residual network with 50 layers, skip connections, pretrained on ImageNet.

**Pros:**

- Transfer learning from ImageNet provides good feature extractors

- Residual connections help with gradient flow
- Robust to vanishing gradient problem

**Cons:**

- **Performance bottleneck:** Only 97.06% accuracy
- **Input mismatch:** Requires upsampling 28x28→224x224 and grayscale→RGB conversion
- **Parameter-heavy:** 24.68M parameters (89.89MB non-trainable)
- Pretrained features may not be optimal for simple grayscale characters

**Why it underperformed:**

1. **Domain mismatch:** ResNet was trained on natural color images (ImageNet), while handwritten characters are simple, high-contrast grayscale patterns
2. **Information loss:** Upscaling 28x28→224x224 creates artificial pixels without adding real information
3. **Overkill architecture:** 50-layer network is excessive for 28x28 character recognition

### 3.3 InceptionV1/GoogLeNet (Pretrained)

**Architecture:** Inception modules with parallel convolutions at different scales.

**Pros:**

- Multi-scale feature extraction
- Efficient parameter usage
- TF Hub integration simplifies implementation

**Cons:**

- **Lowest accuracy:** 95.27% among all models
- Similar input mismatch issues as ResNet
- Transfer learning benefits limited for this specific task

**Why it underperformed:**

Same issues as ResNet (domain mismatch, upscaling artifacts) combined with Inception's complexity being unnecessary for simple character recognition.

### 3.4 MobileNet (Custom Implementation)

**Architecture:** Lightweight CNN using depthwise separable convolutions.

**Pros:**

- **Best performance:** 99.03% accuracy (highest of all models)
- **Optimized for small images:** Input size 64x64 (better than 224x224 for ResNet/Inception)
- **Parameter efficient:** Depthwise separable convolutions reduce computation

- **No domain mismatch:** Trained from scratch on the actual dataset

**Cons:**

- Custom implementation required (not pretrained)
- More complex architecture than VGG
- Requires gradient clipping for stability

**Why it performed best:**

1. **Architecture suitability:** Depthwise separable convolutions are ideal for simple, high-contrast patterns
2. **Input size:** 64x64 is a better fit than 224x224 (ResNet/Inception) or 28x28 (VGG)
3. **Efficiency:** More computation spent on learning relevant features rather than unnecessary complexity
4. **Modern design:** Incorporates BatchNorm, ReLU, and proper regularization

#### 4. Performance Hierarchy

1. **MobileNet (Custom)** → **Best choice** for accuracy and efficiency
2. **VGG19 (Scratch)** → Excellent alternative, simpler implementation
3. **ResNet50V2** → Overkill, not recommended for this task
4. **InceptionV1** → Least suitable for this dataset